

Your next drop is already shot.

From the live laredoute.fr packshots, we dressed 10 real Mango looks on 5 house faces: 50 photoreal on-model images, for a reconstructed ~325 credits of compute (a few dollars), in minutes, with no studio, casting or sample.

Every image was checked by hand against its source packshot. We built it across three brands on your rail, your own label included, so you can judge it on your own product, with a human-QA gate before anything publishes.

Emmanuel de Maistre
Cofounder & CEO, Scenario



10 REAL MANGO LOOKS	5 HOUSE CASTINGS	50 ON-MODEL IMAGES	3 BRANDS INCL. YOUR OWN	~325 CREDITS (RECONSTRUCTED)	minutes NOT WEEKS
----------------------------------	-------------------------------	---------------------------------	---	---	-----------------------------

Built on real garments sourced from laredoute.fr, dressed on five Scenario-designed faces with Pruna P-Image Try-On. The grey-studio backdrop is the verified floor; editorial environments and market localization sit beside it as a labelled capability. One garment is a real La Redoute Collections own-label piece, rendered on the house cast in this deck. Faces are AI-generated.

The bottleneck is the studio, not the photo.

On-model imagery sells. The constraint is never the picture, it is the physical production behind it: book a body, a photographer, a studio, hair, make-up and styling, ship samples, then do it again for the next drop, the next market and the next of thousands of SKUs.

What makes it slow and expensive

- Casting, travel and agency fees to put a range on a represented cast.
- Sample logistics: garments physically present, on the day, in the right size.
- Re-shoots and overruns when a look, a size or a market is missing.
- Scale: a fast-moving assortment, third-party brands included, across 20+ markets.

Why it matters commercially

- On-model context is a known driver of PDP conversion and a known reducer of fit-related returns.
- Representation on the model is increasingly a brand and conversion question, not only an ethical one.
- Speed to a sellable PDP is the gap between a new SKU landing and earning.


We are deliberately not putting invented percentages on those last points. They are well-established directional levers; the pilot in this proposal is designed to measure your real figures, not assert ours.

The rest of this document is the proof, its exact boundary, and the gates that actually decide a deal like this: rights, security and a human-QA step before publish. We built it on your own catalog so you can judge it on your own product.

THE PROOF

What we built on your own catalog, and its exact boundary.

We took one La Redoute brand page (Mango), read 99 live products, selected 10 real looks (6 single pieces plus 4 multi-piece combinations, 15 garment images in all) and dressed 5 Scenario-designed house models in every look. That is 50 on-model images, plus real motion clips and a live own-brand render.



Today on laredoute.fr: the live product pages this was built from (the floral jumpsuit, SKU 602618087). Tomorrow, the same SKU is on-model, in motion, in any market, on any face.

The method, in three steps

STEP	MODEL	WHAT IT DOES
1 · Master	GPT Image 2	Generates one full-body studio house model per casting (~12 credits, ~45s).
2 · Try-on	Pruna P-Image Try-On	Dresses that model in a garment packshot (~3 credits). Up to 11 garment images at once, so a 2 or 3 piece outfit is dressed in a single pass.
3 · Motion	Seedance 2.0 Fast	Animates the approved still into a 6s 9:16 PDP clip, two-keyframe, so the motion is provably the same garment and face.

The boundary, stated plainly. What you see here is one frontal hero per look, one held pose per identity, on a seamless grey backdrop. Back views, multiple angles and alternate poses per model are on the roadmap. Body-size range across the cast is the headline pilot deliverable. They are not implied as done here.

Provenance: the Mango packshots were scrapped from laredoute.fr as a one-time evaluation artifact, quarantined from any production claim. In production, the platform runs only on assets you warrant you have the rights to use. See the rights page.

THE CEILING

The same image, off the grey. In any market.

Studio is the floor, not the ceiling. We took the exact verified try-ons and changed only the world around them: same face, same garment, real light, and the city of your choice. This is the art direction the studio grid cannot show, and the gap between a competent PDP factory and a campaign engine.



Each pair: the verified studio still (left) and the same identity and garment moved into an editorial environment (right).



Localize: the same garment placed in five markets (Paris, Tokyo, New York, Copenhagen, Dubai). Full-body placement uses GPT Image 2 img2img, which holds body proportions; the environment-only swaps use Gemini.

AI capability demonstration. AI-generated environment and AI face. Same identity and garment, preserved from the verified studio still. Not a shipped La Redoute campaign. The studio grid remains the fidelity proof; this is the labelled "what it can look like" layer.

Any brand on the rail. Your own labels included.

La Redoute is a marketplace, not a single label. So we did not stop at Mango. We scraped two more garments live from laredoute.fr, your own-label La Redoute Collections and the tenant brand Vero Moda, and dressed them on the same house cast. The own-brand render is real proof now, not a deferred promise.



Top row: La Redoute Collections "Robe en lin Signature HELOISE" (your own label, SKU 601946493), packshot then on three faces. Bottom row: Vero Moda "Robe a sequins" (SKU 602417412), packshot then on three faces. Each: source packshot left, on-model renders right.

Why this matters. The own-label render is the one image that turns the rights story from a promise into proof: a real garment from your live site, on AI faces, on the same grey-studio fidelity basis as Mango. For marketplace brands, the same pipeline runs the moment the consent path is cleared.

One garment, every face.

Ten looks down, five castings across. A diverse cast (Spanish, East Asian, Black, Northern European, South Asian) for every product, with no casting call, travel or re-shoot. Swap the wearer, keep the garment.



One face, every look. And the pose is a dial.

A reusable, identity-consistent house model is the asset a lookbook or a PDP carousel needs. Here is one face (M5) across all ten Mango looks, the floral jumpsuit included, from a single base shot.



M5 (South Asian, straight-on stance) in all 10 looks. The same face, hair and pose; only the garment changes.

Pose variety, shown not claimed



Same model, same Mango look, three poses generated by re-pose (contrapposto, walking, hand to hair). Alternate poses per model are a pilot deliverable; this is what the capability already produces.

Fidelity, up close. verified

The fair test of try-on is whether it survives a zoom. These were checked by hand against their source packshots.



S2 · floral jumpsuit (the hero look). The large red-and-white print holds its scale across the body, the halter neckline and the waist belt stay intact, and the drape reads as fabric, not a flat overlay.

A colorway win, not a miss verified



Source packshot



On-model render

The C9 bottom is a real burgundy wide-leg trouser, sold under the misleading name *Jean large*. The try-on reproduced that unusual colorway correctly. The only honest note: the source's very wide flare rendered slightly straighter. That is a sourcing and metadata lesson, not a Scenario miss, and a lazier deck would have wrongly confessed it as one.

Why a human-QA gate is mandatory, not optional. Every fidelity finding in this deck was verified against the source packshot at zoom by a named reviewer. The dangerous failure mode in try-on is never visible breakage; it is silent substitution, an attractive photoreal image of the wrong product. The gate exists to catch exactly that before anything publishes, on hard, patterned or cropped SKUs especially.

Rights, IP and synthetic media.

This page is honest about what is built and what is pilot work. The own-brand render now exists on disk and in this deck, so the rights story is a proof, not a promise.

Own-brand rendered

A real La Redoute Collections garment is dressed on the house cast in this deck (see "Every brand"), on assets where the rights are unambiguous. In the pilot we extend it across your own-label lines at volume. The rights story is proof, not a promise.

Third-party brands

The Mango and Vero Moda packshots used here were a one-time evaluation artifact, quarantined from any production claim. In production you feed only assets you warrant you have derivative-image rights to; for marketplace products we set out the brand-consent path.

Synthetic media

These are AI-generated faces. You own the five synthetic identities. We state a directional indemnity stance on IP and likeness for synthetic-media output now, not later, and our position on EU AI Act and French advertising disclosure. Real, contracted talent remains an option.

This is the gate between an interesting demo and a signable pilot. We are clearing it head-on, with a real own-brand render rather than a claim we have not built.

Data and platform posture

- GDPR / DPA posture and EU data residency for La Redoute assets.
- Retention defaults and a no-training-on-your-assets default, with opt-out.
- Per-model sub-processor list and license terms exposed in the Scenario catalog. Production runs only on assets you warrant you have rights to.

Security, data, and how it fits your pipeline.

Scenario is a platform your teams operate, not a consultancy that runs your catalog. It sits between your PIM and your PDP.

PIM / DAM → packshot in → **Scenario** (GPT Image 2 · P-Image Try-On · Seedance 2.0 Fast) → on-model asset keyed to SKU → **PDP**

Integration via API / MCP. Packshot in, SKU-keyed on-model image and clip out, with the human-QA gate before publish.

Service levels (to be set with you)

- Turnaround commitment per batch
- Uptime target
- Support response time
- Rate-lock so rollout pricing is known, not renegotiated

Governance

- EU data residency, GDPR / DPA
- No-training default, retention controls
- Per-model license terms visible in-platform
- Audit trail: every image traceable to its source packshot and reviewer

THE NUMBERS

Unit economics, with the sources on the page.

Demonstrated, hand-verified

50 on-model images across a 5-casting cast for a reconstructed ~325 credits, in minutes. ~3 credits per try-on. A new casting costs ~12 credits (one master) plus ~3 per swap, so representation across faces is a parameter, not a budget line. The ~325 figure is reconstructed from job-log line items (the logs hold job IDs, not a single metered total); the appendix shows it summing.

The cost that actually scales is QA, not compute

Compute is cents. The honest fully-loaded Scenario cost per on-model image is compute plus human art-direction, QA and retouch minutes at the accepted rate. Hard, patterned or cropped SKUs prove that QA labor is non-zero. Even fully loaded, it lands an order of magnitude below a studio image.

The studio benchmark, cited

LINE ITEM	RANGE (PER DAY UNLESS NOTED)
Photographer	\$1,000 to \$2,000 (up to \$5,000+)
Model (agency)	\$500 to \$3,000, plus 20 to 40% agency fee
Studio rental	\$300 to \$2,000
Hair & make-up	\$800 to \$2,400
Styling	\$400 to \$1,000
Retouching	\$20 to \$80 per image

A lean one-day shoot lands near **~\$46 per finished image** (about 60 images for ~\$2,750). A worked two-day, 200-image shoot lands near **~\$96 per image** on line items alone, and the sources note that effective per-image cost runs **2 to 3 times** the quoted rate once coordination, sample logistics and re-shoots are counted. Set against a Scenario image at cents of compute plus minutes of QA, the order-of-magnitude case does not depend on a single number.

Illustrative ROI (labelled illustrative): shoots avoided per category per year from the cost and speed delta, with a sensitivity range on conversion lift and return reduction. The pilot measures your real figures against your own baseline; we do not assert a conversion percentage here.

Sources: squareshot.com (per-image and one-day-shoot benchmark); blendnow.com and nightjar.so (day-rate line items and the 2 to 3x effective-cost finding). Full citations in the appendix. Replace any figure with your own blended cost.

What separates this from a generic AI-imagery pitch.

1 • Built on your live catalog

Real SKUs from laredoute.fr, hard cases included (prints, satin, denim, a three-piece tailored look), across three brands and your own label, not stock demos. You judge it on your own product.

2 • A platform you operate

Your packshots, your identity-locked house-model library, your pipeline via API/MCP, your data governance, with a directional indemnity stance stated upfront. This clears the procurement, IT and EU-compliance bar that pure-service vendors stumble on.

3 • Honesty as the moat

Every number grounded and shown summing, the benchmark cited inline, the motion a real file not a storyboard, the body claim scoped to faces, a colorway conceded as a win, and a real own-brand render rather than a promise.

4 • Stills to video, one platform

GPT Image 2, P-Image Try-On and Seedance 2.0 Fast in one place, proven by a real on-model turn clip. This is infrastructure, not a one-trick demo.

For a buyer who has sat through overclaiming AI demos, proof on your own catalog, your own label included, with every number shown summing, is the differentiator that earns the pilot.

Real clips, not a storyboard.

Five 9:16 on-model clips, one garment (the floral jumpsuit) across all five castings, each generated from an already-approved try-on still via Seedance 2.0 Fast. The clips are two-keyframe interpolations: pose A is the still you inspected, pose B is a re-pose walk of the same garment, so the motion is provably the same garment and the same face, with the garment length pinned at both ends. The drape and hem movement is the part a flat image cannot sell. These cut together into a branded film on the microsite, which also carries one lifestyle beat: the same jumpsuit walking a Paris street, motion grown from the editorial still.

The format

9:16, 720p, ~6s, the everyday PDP and social format. ~176 credits per clip.

How it is made

Two-keyframe mode reuses the approved still as the first frame and a re-pose as the last. No new shoot, no continuity gap, no garment drift.

What it proves

Stills-to-video on one platform is real, today, removing the last unproven capability from the pitch.

The clips ship as real files in this package (media/video) and play in the microsite. A cross-model morph (one garment moving across faces) is an optional capability, deliberately held behind a garment-shape and face-stability QA pass before it represents the brand. Any additional clip volumes in a rollout are projected at the platform rate (~120 to 400 credits per clip), labelled as estimates.



One frame from each of the five real clips: the floral jumpsuit on M1, M2, M3, M4 and M5. One garment, every face, in motion.

It works. Let us prove it on your catalog.

This is not a concept that needs months to validate. Everything in this deck already runs: the master, the try-on, the editorial placement, the localization, the motion, the own-brand render. Scenario has the models, the pipeline and the governance in place today.

What we can show, fast

- Your own SKUs dressed on an identity-locked cast, into a shared review space.
- One look across a body-size range, the dimension that drives fashion returns.
- Localized and editorial variants of the same look, per market.
- Motion from the approved stills, no continuity gap.
- The human-QA gate and acceptance rubric, running on your product.

What you bring

- A slice of live SKUs you want to see on-model, your own label included.
- A casting brief: the faces and the build range that fit your audience.
- An art director to set the acceptance bar.

We run it, you judge it against your own studio baseline. Days, not quarters.

Emmanuel offers a 30-minute working session to pick the SKU slice and lock the success metric.

Cordialement.

Method, costs, provenance, sources.

The looks (live laredoute.fr SKUs)

CODE	LOOK	PRICE
S1	Robe chemisier imprimé floral	39,99 €
S2	Combinaison imprimé floral (jumpsuit, the hero look)	49,99 €
S3	Robe asymétrique imprimé fleuri	79,99 €
S4	Robe satinée asymétrique	69,99 €
S5	Robe bustier en jean	49,99 €
S6	Robe à fleurs cut-out	79,99 €
C7 to C10	Four multi-piece combinations (shirt + trousers, top + skirt, blouse + trousers, shirt + trousers + blazer)	from 72,98 €
+ brands	La Redoute Collections "Robe en lin Signature HELOISE" (own label, SKU 601946493); Vero Moda "Robe a sequins" (SKU 602417412)	live

Reconstructed credit total

ITEM	CREDITS
5 house masters + smoke test (incl. moderation, pose re-fires)	~173
50 try-on swaps (~3 each)	~150
Total (the 50-image demonstrated set)	~323
6 brand / own-label try-ons + re-poses + 5 geo + lifestyle	~174
2 geo redos (GPT Image 2) + music bed	~41
Motion: 6 clips at ~176 (Seedance 2.0 Fast)	~1,056
Whole build, all-in	~1,600

GPT Image 2 medium: ~12 credits, ~45s. Pruna P-Image Try-On: ~3 credits, up to 11 garment references. Seedance 2.0 Fast: ~176 credits at 6s / 720p. A few dollars to low tens for the entire build, on the Public Data project.

Show-your-work, and sources.

Honest gotchas

- GPT Image 2 moderation occasionally re-fired a master; resolved with a more modest base wardrobe. Pose wording matters: "glance over shoulder" produced a back view, corrected to front-facing.
- C9 is a faithful render of a mislabeled SKU (burgundy "Jean large"), with a minor flare-straightening note. A win, not a miss.
- Full-body placement: Gemini stretched the figure on two geo renders; redone with GPT Image 2, which holds proportions. Gemini stays fine for environment-only swaps.
- Motion uses two keyframes (still plus re-pose): real walking motion with the garment pinned, no drift.

Sources (studio cost benchmark)

squareshot.com, "Guide to Product Photography Rates" (per-image and one-day-shoot figures); blendnow.com, "The True Cost of a Professional Photoshoot" (day-rate line items); nightjar.so, "The Real Cost of Product Photography" (effective-cost multiple). Figures are public industry benchmarks, not La Redoute data; replace with your blended cost.

Garments © their respective brands (Mango, Vero Moda) and La Redoute Collections, sourced via laredoute.fr for one-time evaluation, quarantined from production claims. Faces are AI-generated. Built in Scenario on the Public Data project. Prepared for La Redoute by Scenario, Inc.