

Your next Mango drop is already shot.

From the live laredoute.fr packshots, we dressed 10 real Mango looks on 5 house faces: 50 photoreal on-model images, for a reconstructed ~325 credits of compute (a few dollars), in minutes, with no studio, casting or sample.

Every image was checked by hand against its source packshot. The one look our try-on did not yet reproduce faithfully is in this deck too, beside a real on-model video, because the honest version is the one you can put in front of legal.

Emmanuel de Maistre
Cofounder & CEO, Scenario



10	5	50	5	~325	minutes
REAL MANGO LOOKS	HOUSE CASTINGS	ON-MODEL IMAGES	REAL MOTION CLIPS	CREDITS (RECONSTRUCTED)	NOT WEEKS

Built on real Mango garments sourced from laredoute.fr, dressed on five Scenario-designed faces with Pruna P-Image Try-On. Lifestyle and editorial art direction (location, light, styling) are a pilot objective: the seamless studio backdrop here is a starting point, not the ceiling. No own-brand render is claimed on this page; it is a Phase-1 pilot deliverable. Faces are AI-generated.

The bottleneck is the studio, not the photo.

On-model imagery sells. The constraint is never the picture, it is the physical production behind it: book a body, a photographer, a studio, hair, make-up and styling, ship samples, then do it again for the next drop, the next market and the next of thousands of SKUs.

What makes it slow and expensive

- Casting, travel and agency fees to put a range on a represented cast.
- Sample logistics: garments physically present, on the day, in the right size.
- Re-shoots and overruns when a look, a size or a market is missing.
- Scale: a fast-moving assortment, third-party brands included, across 20+ markets.

Why it matters commercially

- On-model context is a known driver of PDP conversion and a known reducer of fit-related returns.
- Representation on the model is increasingly a brand and conversion question, not only an ethical one.
- Speed to a sellable PDP is the gap between a new SKU landing and earning.

We are deliberately not putting invented percentages on those last points. They are well-established directional levers; the pilot in this proposal is designed to measure your real figures, not assert ours.

The rest of this document is the proof, its exact boundary, and the two gates that actually decide a deal like this: rights and accuracy. We built it on your own catalog so you can judge it on your own product.

What we built on your own catalog, and its exact boundary.

We took one La Redoute brand page (Mango), read 99 live products, selected 10 real looks (6 single pieces plus 4 multi-piece combinations, 15 garment images in all) and dressed 5 Scenario-designed house models in every look. That is 50 on-model images, plus 5 real motion clips.

The method, in three steps

STEP	MODEL	WHAT IT DOES
1 · Master	GPT Image 2	Generates one full-body studio house model per casting (~12 credits, ~45s).
2 · Try-on	Pruna P-Image Try-On	Dresses that model in a garment packshot (~3 credits). Up to 11 garment images at once, so a 2 or 3 piece outfit is dressed in a single pass.
3 · Motion	Seedance 2.0 Fast	Animates the approved still into a 5s 9:16 PDP clip (~146 credits), so the motion is provably the same garment and face.

Identity, the model's pose, the shoes and the backdrop are preserved across every swap. Only the clothing changes.

The boundary, stated plainly. What you see here is one frontal hero per look, one held pose per identity, on a seamless grey backdrop, with no own-brand render. Back views, multiple angles, in-context lifestyle scenes, editorial lighting, alternate poses per model and a La Redoute own-brand look are all on the roadmap and in the pilot scope on the closing page. They are not implied as done here.

Provenance: the Mango packshots were scraped from laredoute.fr as a one-time evaluation artifact, quarantined from any production claim. In production, the platform runs only on assets you warrant you have the rights to use. See the rights page.

THE CEILING

The same image, off the grey.

Studio is the floor, not the ceiling. We took the exact verified try-ons and changed only the world around them: same face, same garment, real light. This is the art direction the studio grid cannot show, and the gap between a competent PDP factory and a campaign engine.



Each pair: the verified studio still (left) and the same identity and garment moved into an editorial environment (right). One Gemini pass, ~13 credits each.

AI capability demonstration. AI-generated environment and AI face. Same identity and garment, preserved from the verified studio still. Not a shipped La Redoute campaign. The studio grid remains the fidelity proof; this is the labelled "what it can look like" layer. An own-brand render stays a Phase-1 deliverable.

One garment, every face.

Ten looks down, five castings across. A diverse cast (Spanish, East Asian, Black, Northern European, South Asian) for every product, with no casting call, travel or re-shoot. Swap the wearer, keep the garment.



One face, every look. And the pose is a dial.

A reusable, identity-consistent house model is the asset a lookbook or a PDP carousel needs. Here is one face (M5) across all ten Mango looks, from a single base shot.



M5 (South Asian, straight-on stance) in all 10 looks. The same face, hair and pose; only the garment changes.

Pose variety, shown not claimed



Same model, same Mango look, three poses generated by re-pose (contrapposto, walking, hand to hair). Alternate poses per model are a pilot deliverable; this is what the capability already produces.

Fidelity, up close. verified

The fair test of try-on is whether it survives a zoom. These two were checked by hand against their source packshots.



S4 · red satin. The one-shoulder cowl, the satin sheen and the high-low asymmetric hem all hold.



S2 · floral jumpsuit. The large red-and-white print holds its scale across the body; halter and belt intact.

A colorway win, not a miss verified



Source packshot



On-model render

The C9 bottom is a real burgundy wide-leg trouser, sold under the misleading name *Jean large*. The try-on reproduced that unusual colorway correctly. The only honest note: the source's very wide flare rendered slightly straighter. That is a sourcing and metadata lesson, not a Scenario miss, and a lazier deck would have wrongly confessed it as one.

Where it is not perfect yet. the C10 case

One look did not reproduce faithfully, and it is the most important image in this deck. The honest failure mode is not visible breakage; it is silent substitution.



Source: cropped, short-sleeve, blue-and-white striped shirt.



Render: the shirt became long-sleeve and full-length tucked, and the blue stripe faded to a pale wash.

What held: the brown blazer and the aubergine wide-leg trousers reproduced faithfully. **What was substituted:** the cropped short-sleeve striped shirt rendered as a long-sleeve, full-length tucked shirt, and the crisp blue stripe washed out. The result is an attractive, photoreal image of the wrong product, which is more dangerous than an obviously broken one, because nothing looks wrong.

This is exactly why a human-QA-before-publish gate is mandatory, not optional. Every fidelity finding in this deck, hold or miss, was verified against the source packshot at zoom by a named reviewer. The C10 catch is the proof that the gate runs.

What we propose you adopt

- An art-director acceptance rubric, one line: lighting, styling, drape, print-scale, color-accuracy, silhouette.
- A strawman acceptance threshold: a target share of images publishable unedited, plus a pattern / color / silhouette fidelity pass-rate, agreed before the pilot starts.

Rights, IP and synthetic media.

This page is honest about what is built and what is pilot work. No own-brand proof exists on disk today, so it is named as future everywhere it appears.

Own-brand **Phase 1**

In the pilot we will render one La Redoute Collections garment on-model, on your own assets, where the rights are unambiguous. This makes the rights story a proof, not a promise. We do not claim a shipped own-brand render here, because there is not one yet.

Third-party brands

The Mango packshots used here were a one-time evaluation artifact, quarantined from any production claim. In production you feed only assets you warrant you have derivative-image rights to; for marketplace products we set out the brand-consent path.

Synthetic media

These are AI-generated faces. You own the five synthetic identities. We state a directional indemnity stance on IP and likeness for synthetic-media output now, not later, and our position on EU AI Act and French advertising disclosure. Real, contracted talent remains an option.

This is the gate between an interesting demo and a signable pilot. We are clearing it head-on rather than claiming an own-brand proof we have not built.

Data and platform posture

- GDPR / DPA posture and EU data residency for La Redoute assets.
- Retention defaults and a no-training-on-your-assets default, with opt-out.
- Per-model sub-processor list and license terms exposed in the Scenario catalog. Production runs only on assets you warrant you have rights to.

Security, data, and how it fits your pipeline.

Scenario is a platform your teams operate, not a consultancy that runs your catalog. It sits between your PIM and your PDP.

PIM / DAM → packshot in → **Scenario** (GPT Image 2 · P-Image Try-On · Seedance 2.0 Fast) → on-model asset keyed to SKU → **PDP**

Integration via API / MCP. Packshot in, SKU-keyed on-model image and clip out, with the human-QA gate before publish.

Service levels (to be set in the pilot)

- Turnaround commitment per batch
- Uptime target
- Support response time
- Rate-lock so rollout pricing is known, not renegotiated

Governance

- EU data residency, GDPR / DPA
- No-training default, retention controls
- Per-model license terms visible in-platform
- Audit trail: every image traceable to its source packshot and reviewer

THE NUMBERS

Unit economics, with the sources on the page.

Demonstrated, hand-verified

50 on-model images across a 5-casting cast for a reconstructed ~325 credits, in minutes. ~3 credits per try-on. A new casting costs ~12 credits (one master) plus ~3 per swap, so representation across faces is a parameter, not a budget line. The ~325 figure is reconstructed from job-log line items (the logs hold job IDs, not a single metered total); the appendix shows it summing.

The cost that actually scales is QA, not compute

Compute is cents. The honest fully-loaded Scenario cost per on-model image is compute plus human art-direction, QA and retouch minutes at the accepted rate. The C10 case proves that QA labor on hard, patterned SKUs is non-zero. Even fully loaded, it lands an order of magnitude below a studio image.

The studio benchmark, cited

LINE ITEM	RANGE (PER DAY UNLESS NOTED)
Photographer	\$1,000 to \$2,000 (up to \$5,000+)
Model (agency)	\$500 to \$3,000, plus 20 to 40% agency fee
Studio rental	\$300 to \$2,000
Hair & make-up	\$800 to \$2,400
Styling	\$400 to \$1,000
Retouching	\$20 to \$80 per image

A lean one-day shoot lands near **~\$46 per finished image** (about 60 images for ~\$2,750). A worked two-day, 200-image shoot lands near **~\$96 per image** on line items alone, and the sources note that effective per-image cost runs **2 to 3 times** the quoted rate once coordination, sample logistics and re-shoots are counted. Set against a Scenario image at cents of compute plus minutes of QA, the order-of-magnitude case does not depend on a single number.

Illustrative ROI (labelled illustrative): shoots avoided per category per year from the cost and speed delta, with a sensitivity range on conversion lift and return reduction. The pilot measures your real figures against your own baseline; we do not assert a conversion percentage here.

Sources: squareshot.com (per-image and one-day-shoot benchmark); blendnow.com and nightjar.so (day-rate line items and the 2 to 3x effective-cost finding). Full citations in the appendix. Replace any figure with your own blended cost.

What separates this from a generic AI-imagery pitch.

1 • Built on your live catalog

Real Mango SKUs from laredoute.fr, hard cases included (prints, satin, denim, a three-piece tailored look), not stock demos. You judge it on your own product. An own-brand render is an honest Phase-1 deliverable, named as future.

2 • A platform you operate

Your packshots, your identity-locked house-model library, your pipeline via API/MCP, your data governance, with a directional indemnity stance stated upfront. This clears the procurement, IT and EU-compliance bar that pure-service vendors stumble on.

3 • Honesty as the moat

Every number grounded and shown summing, the benchmark cited inline, the motion a real file not a storyboard, the body claim scoped to faces, the one fidelity miss shown with a confession that matches the pixels, and a colorway conceded as a win rather than mis-confessed.

4 • Stills to video, one platform

GPT Image 2, P-Image Try-On and Seedance 2.0 Fast in one place, proven by a real on-model turn clip. This is infrastructure, not a one-trick demo.

For a buyer who has sat through overclaiming AI demos, a confession that matches the pixels, plus a refusal to manufacture a fake miss, is the differentiator that earns the pilot.

MOTION

Real clips, not a storyboard.

Five 9:16 on-model clips, one per casting, each generated from an already-approved try-on still via Seedance 2.0 Fast. Because the first frame is the still you inspected, the motion is provably the same garment and the same face. The drape and hem movement is the part a flat image cannot sell. These cut together into a 30-second branded film (in the package and on the microsite), which also carries one editorial-environment motion beat: the same red-satin look turning on a dusk terrace, motion grown from the editorial still.

The format

9:16, 720p, ~5s, the everyday PDP and social format. ~146 credits per clip.

How it is made

First-frame mode reuses the approved still. No new shoot, no continuity gap between the photo and the video.

What it proves

Stills-to-video on one platform is real, today, removing the last unproven capability from the pitch.

The five clips ship as real files in this package (media/video) and play in the microsite. A cross-model morph (one garment moving across faces) is an optional capability, deliberately held behind a garment-shape and face-stability QA pass before it represents the brand. Any additional clip volumes in a rollout are projected at the platform rate (~120 to 400 credits per clip), labelled as estimates.



One frame from each of the five real clips: M1 red satin, M2 floral, M3 jumpsuit, M4 linen separates, M5 slip dress.

NEXT STEP

The pilot. Bounded, measurable, signable.

Scope. You pick ~50 to 100 live SKUs including at least one own-brand line (rights unambiguous) plus a fast-moving third-party set, and a casting brief of 3 to 5 identity-locked faces. You art-direct and approve; we provide the platform, onboarding and pipeline templates.

Phase 1 **days, no eng**

Cost and lead-time baseline. Every SKU dressed on the agreed castings, into a shared review space, measured against your studio baseline. Includes the mandatory human-QA gate and the acceptance rubric.

Deliverables that convert claims to evidence:

- one La Redoute own-brand look rendered on your assets;
- one look across XS-to-XXL builds (body-size, the returns lever);
- 2 to 3 alternate poses on one model;
- 2 to 3 localized variants.

Phase 2 **weeks, optional**

A controlled PDP A/B on a SKU subset, measuring conversion and fit-related return rate against existing imagery, with minimum sample size, duration and pass thresholds agreed up front. Requires your engineering.

Acceptance criteria

Art-director acceptance rate (share publishable unedited) against the rubric, a pattern/color/silhouette fidelity pass-rate, turnaround, and fully-loaded cost per look.

Commercials, decided not blank

A fixed pilot price, or free with the reason stated; a directional platform-license and annual-commit range for rollout; an estimated credit consumption for the agreed SKU x casting x motion volume; asset ownership by La Redoute; usage and IP terms for third-party-derived imagery; a directional indemnity stance; named SLA values; and a rate-lock so the Phase-2 price is known. **Decision gate:** a dated go / no-go with the rollout price and the single number the pilot was built to produce.

Emmanuel offers a 30-minute working session to pick the SKU slice and lock the success metric.

Cordialement.

APPENDIX

Method, costs, provenance, sources.

The looks (live laredoute.fr SKUs)

CODE	LOOK	PRICE
S1	Robe chemisier imprimé floral	39,99 €
S2	Combinaison imprimé floral (jumpsuit)	49,99 €
S3	Robe asymétrique imprimé fleuri	79,99 €
S4	Robe satinée asymétrique	69,99 €
S5	Robe bustier en jean	49,99 €
S6	Robe à fleurs cut-out	79,99 €
C7 to C10	Four multi-piece combinations (shirt + trousers, top + skirt, blouse + trousers, shirt + trousers + blazer)	from 72,98 €

Reconstructed credit total (~325)

ITEM	CREDITS
Smoke test (master variants + 1 swap + 1 re-pose)	~77
5 house masters (incl. moderation and pose re-fires)	~96
50 try-on swaps (~3 each)	~150
Total (stills)	~323
Motion: 5 clips at ~146 (separate capability)	~730

GPT Image 2 medium: ~12 credits, ~45s. Pruna P-Image Try-On: ~3 credits, up to 11 garment references.

Seedance 2.0 Fast: ~146 credits at 5s / 720p.

Honest gotchas (the show-your-work)

- GPT Image 2 output moderation occasionally re-fired a master; resolved with a more modest base wardrobe.
- Pose wording matters: "glance over shoulder" produced a back view; corrected to front-facing.
- C9 is a faithful render of a mislabeled SKU (burgundy "Jean large"), with a minor flare-straightening note. A win, not a miss.
- C10 is a silhouette substitution: short-to-long sleeve, crop-to-full-length tucked, stripe faded. The reason the QA gate is mandatory.

Sources (studio cost benchmark)

squareshot.com, "Guide to Product Photography Rates" (per-image and one-day-shoot figures); blendnow.com, "The True Cost of a Professional Photoshoot" (day-rate line items); nightjar.so, "The Real Cost of Product Photography" (effective-cost multiple). Figures are public industry benchmarks, not La Redoute data; replace with your blended cost.